

Final: STAT 570

Philip Pham

December 9, 2018

Consider the failure time data in Table 1.

1. We describe a simple model for these data. Let p ($0 < p < 1$) denote the weekly failure probability, i.e., the probability of failure during any week, and T the random variable describing the week at which failure occurred. Then T may be modeled as a geometric random variable:

$$\mathbb{P}(T = t | p) = \begin{cases} p(1-p)^{t-1}, & t = 1, 2, \dots; \\ 0, & \text{otherwise.} \end{cases} \quad (1)$$

Let Y_t represent the number of components that fail in week t , $t = 1, 2, \dots, N$, and Y_{N+1} the number of components that have not failed by week N .

- (a) Show that the likelihood function is

$$L(p) = \left[(1-p)^N\right]^{Y_{N+1}} \prod_{t=1}^N \left[p(1-p)^{t-1}\right]^{Y_t}. \quad (2)$$

Solution: An individual component's failure week has distribution Geometric(p). The probability that a single component fails in week t is the probability that it survived $t-1$ weeks and failed on week t , which is $p(1-p)^{t-1}$ from Equation 1. There are Y_t such components, which gives us the factors for $t = 1, 2, \dots, N$.

The probability that a component fails at a later date is

$$(1-p)^N \sum_{k=1}^{\infty} p(1-p)^{k-1} = (1-p)^N \frac{p}{1-(1-p)} = (1-p)^N,$$

which gives us the remaining factor. There are Y_{N+1} remaining components, so

$$L(p) = \left\{ \prod_{t=1}^N \left[p(1-p)^{t-1}\right]^{Y_t} \right\} \times \left[(1-p)^N\right]^{Y_{N+1}}.$$

- (b) Find an expression for the MLE \hat{p} .

Solution: The score function is

$$\begin{aligned}
S(p) &= \frac{\partial}{\partial p} \log L(p) \\
&= \frac{\partial}{\partial p} \left[NY_{N+1} \log(1-p) + \sum_{t=1}^N Y_t (\log p + (t-1) \log(1-p)) \right] \\
&= -\frac{NY_{N+1}}{1-p} + \sum_{t=1}^N Y_t \left(\frac{1}{p} - \frac{t-1}{1-p} \right) = -\frac{NY_{N+1}}{1-p} + \sum_{t=1}^N Y_t \frac{1-pt}{p(1-p)}. \quad (3)
\end{aligned}$$

Solving for $S(\hat{p}) = 0$, we find the MLE:

$$\hat{p} \left(NY_{N+1} + \sum_{t=1}^N tY_t \right) = \sum_{t=1}^N Y_t \implies \boxed{\hat{p} = \frac{\sum_{t=1}^N Y_t}{NY_{N+1} + \sum_{t=1}^N tY_t}}. \quad (4)$$

- (c) Find the form of the observed information and hence the asymptotic variance of the maximum likelihood estimate (MLE).

Solution: Using Equation 3, the expected observed information is

$$\begin{aligned}
I(p) &= \mathbb{E} \left[-\frac{\partial}{\partial p} S(p) \mid p \right] \\
&= \frac{N\mathbb{E}[Y_{N+1} \mid p]}{(1-p)^2} + \sum_{t=1}^N \mathbb{E}[Y_t \mid p] \left(\frac{1}{p^2} + \frac{t-1}{(1-p)^2} \right) \\
&= n \frac{(1-p)^N}{(1-p)^2} + np \sum_{t=1}^N (1-p)^{t-1} \left(\frac{1}{p^2} + \frac{t-1}{(1-p)^2} \right) \\
&= n \left[\frac{(1-p)^N}{(1-p)^2} + \frac{1 - (1-p)^N}{p^2} + \frac{(1-p) - (1-p)^N}{p(1-p)^2} \right] \\
&= \boxed{n \frac{1 - (1-p)^N}{p^2(1-p)}}, \quad (5)
\end{aligned}$$

where $n = Y_{N+1} + \sum_{t=1}^N Y_t$.

From Equation 5, the asymptotic variance of \hat{p} is

$$\text{var}(\hat{p}) \approx \text{vâr}(\hat{p}) = I(\hat{p})^{-1} = \frac{1}{n} \times \frac{\hat{p}^2(1-\hat{p})}{1-(1-\hat{p})^N} \quad (6)$$

by asymptotic normality of the MLE.

- (d) For the data in Table 1, calculate the MLE, \hat{p} , the variance of \hat{p} , and an asymptotic 95% confidence interval for p .

Solution: The MLE can be calculated with Equation 4 to be $\boxed{\hat{p} = 0.354717}$.

The variance can be found with Equation 6 to be $\boxed{\text{vâr}(\hat{p}) = 0.00016828}$.

If Φ is the cumulative distribution function for a standard normal, we can use asymptotic normality to find the 95% confidence interval as

$$\left[\hat{p} + \Phi^{-1}(0.025) \sqrt{\text{vâr}(\hat{p})}, \hat{p} + \Phi^{-1}(0.975) \sqrt{\text{vâr}(\hat{p})} \right] = \boxed{[0.32929, 0.38014]}.$$

- (e) We now consider a Bayesian analysis. The conjugate prior for p is a beta distribution, $\text{Beta}(a, b)$. State the form of the posterior with this choice. Give the form of the posterior mean and write as a weighted combination of the MLE and the prior mean.

Solution: By Bayes' rule, we know the posterior density is proportional to the likelihood times the prior. From Equation 2, we'll have

$$\begin{aligned} L(p) \times [p^{a-1} (1-p)^{b-1}] &= p^{a-1} (1-p)^{b+NY_{N+1}-1} \prod_{t=1}^N [p(1-p)^{t-1}]^{Y_t} \\ &= p^{a+\sum_{t=1}^N Y_t-1} (1-p)^{b+\sum_{t=1}^N (t-1)Y_t+NY_{N+1}-1}, \end{aligned}$$

whose form we recognize as the integrand of beta function, so the posterior also has beta distribution, that is,

$$\begin{aligned} p \mid Y_1, Y_2, \dots, Y_{N+1} &\sim \text{Beta} \left(a + \sum_{t=1}^N Y_t, b + \sum_{t=1}^N (t-1)Y_t + NY_{N+1} \right) \\ &= \frac{\Gamma(a') \Gamma(b')}{\Gamma(a'+b')} p^{a'-1} (1-p)^{b'-1}, \end{aligned} \quad (7)$$

where $a' = a + \sum_{t=1}^N Y_t$ and $b' = b + \sum_{t=1}^N (t-1)Y_t + NY_{N+1}$.

The posterior mean takes the form

$$\begin{aligned} \mathbb{E}[p \mid Y_1, Y_2, \dots, Y_{N+1}] &= \frac{a'}{a'+b'} \\ &= \frac{a + \sum_{t=1}^N Y_t}{a + b + \sum_{t=1}^N tY_t + NY_{N+1}}. \end{aligned} \quad (8)$$

We have that the prior mean is $p_{\text{prior}} = \frac{a}{a+b}$. Equation 8 can be rewritten as

$$\boxed{\frac{(a+b)p_{\text{prior}} + \left(\sum_{t=1}^N tY_t + NY_{N+1}\right)\hat{p}}{a+b+\sum_{t=1}^N tY_t + NY_{N+1}}}, \quad (9)$$

so the posterior mean is a convex combination of the prior mean and MLE.

- (f) Suppose we wish to fix the parameters of the prior, a and b , so that the mean is μ and the prior standard deviation is σ . Obtain expressions for a and b in terms of μ and σ^2 .

Solution: It is well known that the mean and variance of the $\text{Beta}(a, b)$ distribution are $\frac{a}{a+b}$ and $\frac{ab}{(a+b)^2(a+b+1)}$, respectively. Solving equations

$$\begin{aligned} \frac{a}{a+b} &= \mu \\ \frac{ab}{(a+b)^2(a+b+1)} &= \sigma^2, \end{aligned}$$

we find that

$$a = \mu \left[\frac{\mu(1-\mu)}{\sigma^2} - 1 \right] \quad (10)$$

$$b = (1-\mu) \left[\frac{\mu(1-\mu)}{\sigma^2} - 1 \right]. \quad (11)$$




Figure 1: 2,048 samples drawn from the posterior in Equation 13. The red ticks denote the 5%, 50% and 95% quantiles.

- (g) For the data in Table 1, assume we wish to have a beta prior with $\mu = 0.2$ and $\sigma = 0.08$. State the posterior for the prior corresponding to this choice and evaluate the posterior mean. Simulate samples from the posterior distribution. Provide a histogram representation of the posterior distribution and calculate the 5%, 50% and 95% points of the posterior distribution.

Solution: Apply Equations 10 and 11 with $\mu = 0.2$ and $\sigma = 0.08$, we find the prior:

$$p \sim \text{Beta}(4.8, 19.2). \quad (12)$$

Using Equation 7, we have find the posterior:

$$p \sim \text{Beta}(474.8, 874.2). \quad (13)$$

A histogram of samples drawn from the distribution in Equation 13 can be found in Figure 1. The 5%, 50%, and 95% posterior quantiles are 0.33070873, 0.35189124, and 0.37346975, respectively.

Code for the histogram can be found in `failure_time.ipynb`.

2. (a) A more complex likelihood for these data would assume that the i -th component had their own probability p_i , with the p_i 's arising from a distribution $\pi(p)$. Show that

$$\mathbb{P}(T = t) = \mathbb{E}[(1 - p)^{t-1}] - \mathbb{E}[(1 - p)^t], \quad (14)$$

and

$$\mathbb{P}(T > t) = \mathbb{E}[(1 - p)^t]. \quad (15)$$

Solution: First let us find the survival function in 15.

$$\begin{aligned}
\mathbb{P}(T > t) &= \int_0^1 \mathbb{P}(T > t \mid p) \pi(p) \, dp = \int_0^1 \left[\sum_{s=t+1}^{\infty} p(1-p)^{s-1} \right] \pi(p) \, dp \\
&= \int_0^1 \left[p \sum_{s=0}^{\infty} (1-p)^s \right] (1-p)^t \pi(p) \, dp \\
&= \int_0^1 \left[p \times \frac{1}{1-(1-p)} \right] (1-p)^t \pi(p) \, dp = \int_0^1 (1-p)^t \pi(p) \, dp \\
&= \mathbb{E}[(1-p)^t],
\end{aligned}$$

which proves Equation 15.

The probability mass function in Equation 14 follows:

$$\mathbb{P}(T = t) = \mathbb{P}(T > t-1) - \mathbb{P}(T > t) = \mathbb{E}[(1-p)^{t-1}] - \mathbb{E}[(1-p)^t].$$

- (b) Obtain expressions for $\mathbb{P}(T = t \mid \alpha, \beta)$ and $\mathbb{P}(T > t \mid \alpha, \beta)$ with $\pi(\cdot)$ taken as the beta distribution, $\text{Beta}(\alpha, \beta)$.

Solution: These follow from Equations 14 and 15.

$$\begin{aligned}
\mathbb{P}(T > t) &= \mathbb{E}[(1-p)^t] = \sum_{s=t}^{\infty} \mathbb{E}[p(1-p)^s] \tag{16} \\
&= \sum_{s=t}^{\infty} \int_0^p p(1-p)^s \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} p^{\alpha-1} (1-p)^{\beta-1} \, dp \\
&= \sum_{s=t}^{\infty} \int_0^p \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} p^{\alpha+1-1} (1-p)^{\beta+s-1} \, dp \\
&= \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} \sum_{s=t}^{\infty} \frac{\Gamma(\alpha+1)\Gamma(\beta+s)}{\Gamma(\alpha+\beta+s+1)} = \alpha \frac{\Gamma(\alpha+\beta)}{\Gamma(\beta)} \sum_{s=t}^{\infty} \frac{\Gamma(\beta+s)}{\Gamma(\alpha+\beta+s+1)} \\
&= 1 - \alpha \frac{\Gamma(\alpha+\beta)}{\Gamma(\beta)} \sum_{s=0}^{t-1} \frac{\Gamma(\beta+s)}{\Gamma(\alpha+\beta+s+1)} \\
&= 1 - \frac{1}{B(\alpha, \beta)} \sum_{s=0}^{t-1} B(\alpha+1, \beta+s) = \frac{B(\alpha, \beta+t)}{B(\alpha, \beta)} = \frac{\Gamma(\alpha+\beta)\Gamma(\beta+t)}{\Gamma(\beta)\Gamma(\alpha+\beta+t)}
\end{aligned}$$

where B is the beta function, and we know $\mathbb{P}(T > 0) = 1$.

Plugging Equation 16 into Equation 14, one obtains

$$\mathbb{P}(T = t) = \frac{B(\alpha+1, \beta+t-1)}{B(\alpha, \beta)} = \alpha \frac{\Gamma(\alpha+\beta)\Gamma(\beta+t-1)}{\Gamma(\beta)\Gamma(\alpha+\beta+t)} \tag{17}$$

for $t \in \mathbb{N}$.

- (c) Using the previous part, write down the likelihood function $L(\alpha, \beta)$ corresponding to data $\{Y_t\}_{t=1}^{N+1}$.

Solution: Our model for T is different, so we can substitute Equations 17 and 16 into Equation 2: we'll have $\mathbb{P}(T = t)$ in place of $p(1-p)^{t-1}$ and $\mathbb{P}(T > N)$

in place of $(1 - p)^N$.

$$\begin{aligned} L(\alpha, \beta) &= [\mathbb{P}(T > N)]^{Y_{N+1}} \prod_{t=1}^N [\mathbb{P}(T = t)]^{Y_t} \\ &= \left[\frac{B(\alpha, \beta + N)}{B(\alpha, \beta)} \right]^{Y_{N+1}} \prod_{t=1}^N \left[\frac{B(\alpha + 1, \beta + t - 1)}{B(\alpha, \beta)} \right]^{Y_t}. \end{aligned} \quad (18)$$

(d) Find the MLEs $\hat{\alpha}$ and $\hat{\beta}$ for the data of Table 1.

Solution: From Equation 18, we can consider the log-likelihood function:

$$\begin{aligned} l(\alpha, \beta) &= \log L(\alpha, \beta) \\ &= -n \log B(\alpha, \beta) + Y_{N+1} \log B(\alpha, \beta + N) + \sum_{t=1}^N Y_t \log B(\alpha + 1, \beta + t - 1). \end{aligned} \quad (19)$$

The score function is

$$\begin{aligned} S(\alpha, \beta) &= \nabla l(\alpha, \beta) = \begin{pmatrix} S_\alpha(\alpha, \beta) \\ S_\beta(\alpha, \beta) \end{pmatrix} \\ S_\alpha(\alpha, \beta) &= -n [\psi(\alpha) - \psi(\alpha + \beta)] + Y_{N+1} [\psi(\alpha) - \psi(\alpha + \beta + N)] \\ &\quad + \sum_{t=1}^N Y_t [\psi(\alpha + 1) - \psi(\alpha + \beta + t)], \\ S_\beta(\alpha, \beta) &= -n [\psi(\beta) - \psi(\alpha + \beta)] + Y_{N+1} [\psi(\beta + N) - \psi(\alpha + \beta + N)] \\ &\quad + \sum_{t=1}^N Y_t [\psi(\beta + t - 1) - \psi(\alpha + \beta + t)], \end{aligned} \quad (20)$$

where $\psi(x) = \Gamma'(x) / \Gamma(x)$ is the digamma function.

Numerically solving Equation 20 for $S(\hat{\alpha}, \hat{\beta}) = \mathbf{0}$, I obtain $\hat{\alpha} = 1.413336$

and $\hat{\beta} = 1.38001102$ for the MLEs.

3. (a) Show that the likelihood in Equation 2 can be written as a product of binomial distributions.

Solution: We can model the data as taking N draws from a binomial distribution. Following each draw, we discard the failures and make another draw if $t < N$:

$$\begin{aligned} L(p) &= \prod_{t=1}^N \left[\binom{n - \sum_{s=1}^{t-1} Y_s}{Y_t} p^{Y_t} (1 - p)^{n - \sum_{s=1}^t Y_s} \right] \\ &= \prod_{t=1}^N \left[\binom{\sum_{s=t}^{N+1} Y_s}{Y_t} p^{Y_t} (1 - p)^{\sum_{s=t+1}^{N+1} Y_s} \right], \end{aligned} \quad (21)$$

which is equivalent to Equation 2 up to a constant of proportionality,

In Equation 21, we have a product of binomial probability mass functions, where $Y_t | Y_1, \dots, Y_{t-1} \sim \text{Binomial}(n - \sum_{s=1}^{t-1} Y_s, p)$.

- (b) Fit the binomial model, and show that the estimate of the probability is identical to that under the previous MLE analysis. Obtain a 95% asymptotic confidence interval for p .

Solution: Since Equation 21 only differs from Equation 2 by a constant of proportionality, the score function is also Equation 3. Thus, the MLE is same $\hat{p} = 0.354717$.

The observed information will also be the same. To calculate the expected observed information, we can use the law of total expectation and strong induction. For the base case $\mathbb{E}[Y_1] = np$. In general, $\mathbb{E}[Y_t] = np(1-p)^{t-1}$ for $t = 1, 2, \dots, N$. For $t > 1$, we have

$$\begin{aligned}\mathbb{E}[Y_t] &= \mathbb{E}[\mathbb{E}[Y_t | Y_1, \dots, Y_{t-1}]] = \mathbb{E}\left[p\left(n - \sum_{s=1}^{t-1} Y_s\right)\right] \\ &= p\left(n - \sum_{s=1}^{t-1} \mathbb{E}[Y_s]\right) = p\left(n - \sum_{s=1}^{t-1} np(1-p)^{s-1}\right) \\ &= np\left(1 - p \sum_{s=0}^{t-2} (1-p)^s\right) = np\left(1 - p \frac{1 - (1-p)^{t-1}}{p}\right) \\ &= np(1-p)^{t-1},\end{aligned}\tag{22}$$

which is same as it was under the geometric model.

For Y_{N+1} , we have

$$\begin{aligned}\mathbb{E}[Y_{N+1}] &= \mathbb{E}[\mathbb{E}[Y_{N+1} | Y_1, Y_2, \dots, Y_N]] = \mathbb{E}\left[n - \sum_{t=1}^N Y_t\right] \\ &= n - \sum_{t=1}^N \mathbb{E}[Y_t] = n - np \sum_{t=1}^N (1-p)^{t-1} \\ &= n - np \frac{1 - (1-p)^N}{p} = (1-p)^N,\end{aligned}\tag{23}$$

which is also the same as under the geometric model. Therefore, the expected observed information is the same as Equation 5.

Then, the asymptotic 95% confidence interval for p is also $[0.32929, 0.38014]$.

- (c) Obtain Pearson residuals and comment on the fit of the model.

Solution: Define $n_t = n - \sum_{s=1}^{t-1} Y_s$. The Pearson residuals are

$$\epsilon_t^* = \frac{Y_t - n_t \hat{p}}{\sqrt{n_t \hat{p} (1 - \hat{p})}}\tag{24}$$

for $t = 1, 2, \dots, N$.

The residuals are plotted in Figure 2. Clearly, there are not independent with respect to time. For earlier time steps, the probability of failure is underestimated, and for later time steps, the probability of failure is overestimated. Thus, component failure may not be a memoryless process as assumed by our model.

The Q-Q plot in Figure 3 indicates that the sample residuals are overdispersed relative to the theoretical quantiles. Since only the asymptotic behavior is normal, given only 12 observations, we might expect some deviation from normality, but qualitatively, the deviation is quite significant.

- (d) Fit a binomial model you feel is appropriate.




Figure 2: Residuals as a function of time with the MLE estimate.




Figure 3: Q-Q plot of residuals for the binomial model.

	MLE	Standard error	95% CI lower bound	95% CI upper bound
$\hat{\beta}_0$	-0.090505	0.091144	-0.269143	0.088133
$\hat{\beta}_1$	-0.174740	0.026138	-0.225969	-0.123511

Table 1: MLE estimates for the model in Equation 25.

Solution: From Figure 2, there seems to be some sort of time dependence on the failure probability. That is, the probability of failure is not memoryless. Intuitively, one might hypothesize the longer that a component has survived, the more robust it is, so the probability of failure should decrease with time. To test this hypothesis, we might fit the model:

$$Y_t \mid Y_1, Y_2, \dots, Y_{t-1} \sim \text{Binomial}(n_t, p_t) \quad (25)$$

$$\text{logit}(p_t) = \beta_0 + \beta_1 t.$$

The score function for such a model would be obtained by replacing p with p_t in Equation 21, taking the log, and deriving with respect to $\beta = (\beta_0 \ \beta_1)^\top$.

$$S(\beta) = \sum_{t=1}^N (Y_t - n_t p_t) \begin{pmatrix} 1 \\ t \end{pmatrix}. \quad (26)$$

The observed information is then

$$J(\beta) = \sum_{t=1}^N n_t p_t (1 - p_t) \begin{pmatrix} 1 & t \\ t & t^2 \end{pmatrix}. \quad (27)$$

Using Equations 26 and 27, we obtain MLE estimates and confidence intervals in Table 2. The 95% confidence interval for β_1 does not contain time, so we could reject a null hypothesis $\beta_1 = 0$ in test with significance level 0.05. So, there is likely some correlation between time and failure probability.

Figures 4 and 5 indicate that our new model is much more appropriate. In Figure 4, there's no longer any obvious relationship between time and the residuals. In Figure 5, we see the distribution of the residuals is quite close to normal.

Code for histograms and model fitting can be found in `failure_time.ipynb`.




Figure 4: Residuals for the model in Equation 25 that lets the probability depend on time.




Figure 5: Q-Q plot for the residuals plotted in Figure 4.

Time (weeks), i	Failures, y_i	Temperature, x_i
1	210	24.0
2	108	26.0
3	58	24.0
4	40	26.0
5	17	25.0
6	10	22.0
7	7	23.0
8	6	20.0
9	5	21.0
10	4	18.0
11	2	17.0
12	3	20.0
> 12	15	

Table 2: Time until failure for $n = 485$ components, along with average weekly temperature.